

```

1.  """
2.  Created on Mon Mar 01 14:47:27 2024
3.  @author: Bing
4.  The core code of the LBDSAC algorithm
5.  """
6.  import os
7.  os.environ["CUDA_VISIBLE_DEVICES"] = '0'
8.  import tensorflow as tf
9.  import shutil
10. import numpy as np
11. import pandas as pd
12. import time
13. import datetime
14. import seaborn as sns
15. from logx import EpochLogger
16. from mpi_tf import sync_all_params, MpiAdamOptimizer
17. from mpi_tools import mpi_fork, mpi_sum, proc_id, mpi_statistics_scalar, num_procs
18. from scipy.stats import norm
19. import globalvar as gl
20. import json
21. import argparse
22. from run_utils import setup_logger_kwargs
23. import save_data
24. import matplotlib.pyplot as plt
25. from env2 import unPower
26. from model_modules import placeholder, placeholders, cnn_mlp, get_vars, count_vars, gaussian_likelihood, \
27.     get_target_update, mlp_gaussian_policy, apply_squashing_func, mlp_actor, mlp_critic, Fig
28. import ReplayBuffer
29.
30. n_T = gl.get_value('n_T')
31. n_wind = gl.get_value('n_wind')
32. n_load = gl.get_value('n_load')
33. n_gen = gl.get_value('n_gen')
34. n_wind_data = gl.get_value('n_wind_data')
35. obs_dim = gl.get_value('state_dim')
36. act_dim = gl.get_value('action_dim')
37. P_load = gl.get_value('P_load')
38. P_wind = gl.get_value('P_wind')
39. load_capacity = gl.get_value('load_capacity')
40. wind_capacity = gl.get_value('wind_capacity')
41. EPS = 1e-8
42.
43. snow = datetime.datetime.now()
44. seed_now = str(snow.month)+str(snow.day)+str(snow.hour)

```

```

45. seed_now = int(seed_now)
46.
47. parser = argparse.ArgumentParser()
48. parser.add_argument('--Train', type=bool, default=True, help='True or False')
49. parser.add_argument('--Continue', type=bool, default=False, help='True or False')
50. parser.add_argument('--Portion', type=bool, default=False, help='True or False')
51. parser.add_argument('--SaveFig', type=bool, default=False, help='True or False')
52. parser.add_argument('--load_buffer', type=bool, default=False, help='True or False')
53. parser.add_argument('--Model_version', type=int, default=8000, help='8000')
54. parser.add_argument('--epochs', type=int, default=8000, help='800')
55. parser.add_argument('--steps_per_epoch', type=int, default=20, help='20')
56. parser.add_argument('--update_freq', type=int, default=720, help='720')
57. parser.add_argument('--save_freq', type=int, default=1000, help='500')
58. parser.add_argument('--cost_lim', type=float, default=1e-3, help='1e-3')
59. parser.add_argument('--lr', type=float, default=1e-3, help='1e-3')
60. parser.add_argument('--hidden_sizes_actor', type=list, default=[128, 256, 256, 64], help='actor')
61. parser.add_argument('--hidden_sizes_critic', type=list, default=[128, 256, 64], help='critic')
62. parser.add_argument('--hidden_sizes_var', type=list, default=[128, 256, 64], help='var')
63. parser.add_argument('--cnn_actor', type=list, default=[], help='CNNs of policy network actor')
64. parser.add_argument('--cnn_critic', type=list, default=[], help='CNNs of network critic')
65. parser.add_argument('--cnn_var', type=list, default=[], help='CNNs of network var')
66. parser.add_argument('--hid', type=int, default=512)
67. parser.add_argument('--l', type=int, default=2)
68. parser.add_argument('--gamma', type=float, default=0.99)
69. parser.add_argument('--seed', '-s', type=int, default=seed_now)
70. parser.add_argument('--exp_name', type=str, default='sac')
71. parser.add_argument('--cpu', type=int, default=1)
72. parser.add_argument('--render', default=False, action='store_true')
73. parser.add_argument('--local_start_steps', default=3600, type=int, help='3600')
74. parser.add_argument('--local_update_after', default=3600, type=int, help='3600')
75. parser.add_argument('--batch_size', default=256, type=int) # 256
76. parser.add_argument('--fixed_entropy_bonus', default=None, type=float)
77. parser.add_argument('--entropy_constraint', type=float, default=-1)
78. parser.add_argument('--fixed_cost_penalty', default=None, type=float)
79. parser.add_argument('--cost_constraint', type=float, default=None)
80. parser.add_argument('--lr_s', type=int, default=50)
81. parser.add_argument('--damp_s', type=int, default=10)
82. parser.add_argument('--logger_kwargs_str', type=json.loads, default={'"output_dir": "/.logger"})')
83. args = parser.parse_args()
84.
85. """
86. Soft Actor-Critic with Logic-Based Benders Decomposition Algorithm
87. """
88.

```

```

89. def LBDSAC(env_fn, actor_fn=mlp_actor, critic_fn=mlp_critic,
90.            ac_kwargs_actor=dict(), ac_kwargs_critic=dict(), ac_kwargs_var=dict(),
91.            seed=0, steps_per_epoch=1200, epochs=101, replay_size=500 * 2 * 720, gamma=0.99, cl=0.5,
92.            polyak=0.995, lr=1e-4, batch_size=1024, local_start_steps=600,
93.            max_ep_len=n_T, logger_kwargs=dict(), save_freq=50, local_update_after=int(1e3),
94.            update_freq=120, render=False,
95.            fixed_entropy_bonus=None, entropy_constraint=-1.0,
96.            fixed_cost_penalty=None, cost_constraint=None, cost_lim=None,
97.            reward_scale=1, lr_scale=1, damp_scale=0, Train=True, Continue_Training=False,
98.            Model_Portion=False, load_buffer=True
99.            ):
100.
101.    use_costs = fixed_cost_penalty or cost_constraint or cost_lim
102.    logger = EpochLogger(**logger_kwargs)
103.    logger.save_config(locals())
104.
105.    # Env instantiation
106.    env, test_env = env_fn(), env_fn()
107.
108.    #Setting seeds
109.    seed += 200 * proc_id()
110.    tf.set_random_seed(seed)
111.    np.random.seed(seed)
112.
113.    # Inputs to computation graph
114.    x_ph, a_ph, x2_ph, r_ph, d_ph, c1_ph, c2_ph = placeholders(obs_dim, act_dim, obs_dim, None, None, None, None
115.    )
116.
117.    # Main outputs from computation graph
118.    with tf.variable_scope('main'):
119.        mu, pi, logp_pi = actor_fn(x_ph, a_ph, **ac_kwargs_actor)
120.        qr1, qr1_pi = critic_fn(x_ph, a_ph, pi, name='qr1', **ac_kwargs_critic)
121.        qr2, qr2_pi = critic_fn(x_ph, a_ph, pi, name='qr2', **ac_kwargs_critic)
122.        qc1, qc1_pi = critic_fn(x_ph, a_ph, pi, name='qc1', **ac_kwargs_critic)
123.        qc2, qc2_pi = critic_fn(x_ph, a_ph, pi, name='qc2', **ac_kwargs_critic)
124.
125.    with tf.variable_scope('main', reuse=True):
126.        # Additional policy output from a different observation placeholder
127.        _, pi2, logp_pi2 = actor_fn(x2_ph, a_ph, **ac_kwargs_actor)
128.
129.    # Target value network
130.    with tf.variable_scope('target'):
131.        _, qr1_pi_targ = critic_fn(x2_ph, a_ph, pi2, name='qr1', **ac_kwargs_critic)
132.        _, qr2_pi_targ = critic_fn(x2_ph, a_ph, pi2, name='qr2', **ac_kwargs_critic)
133.        _, qc1_pi_targ = critic_fn(x2_ph, a_ph, pi2, name='qc1', **ac_kwargs_critic)

```

```

132.     _, qc2_pi_targ = critic_fn(x2_ph, a_ph, pi2, name='qc2', **ac_kwargs_critic)
133.
134.     # Entropy bonus
135.     if fixed_entropy_bonus is None:
136.         with tf.variable_scope('entreg'):
137.             soft_alpha = tf.get_variable('soft_alpha',
138.                                         initializer=0.0,
139.                                         trainable=True,
140.                                         dtype=tf.float32)
141.             alpha = tf.nn.softplus(soft_alpha)
142.     else:
143.         alpha = tf.constant(fixed_entropy_bonus)
144.     log_alpha = tf.log(tf.clip_by_value(alpha, 1e-8, 1e8))
145.
146.     # Cost penalty
147.     if use_costs:
148.         if fixed_cost_penalty is None:
149.             with tf.variable_scope('costpen1'):
150.                 soft_beta1 = tf.get_variable('soft_beta1',
151.                                             initializer=0.0,
152.                                             trainable=True,
153.                                             dtype=tf.float32)
154.                 beta1 = tf.nn.softplus(soft_beta1)
155.                 log_beta1 = tf.log(tf.clip_by_value(beta1, 1e-8, 1e8))
156.         else:
157.             beta1 = tf.constant(fixed_cost_penalty)
158.             log_beta1 = tf.log(tf.clip_by_value(beta1, 1e-8, 1e8))
159.     else:
160.         beta1 = 0.0 # costs do not contribute to policy optimization
161.     print('Not using costs')
162.
163.     if use_costs:
164.         if fixed_cost_penalty is None:
165.             with tf.variable_scope('costpen2'):
166.                 soft_beta2 = tf.get_variable('soft_beta2',
167.                                             initializer=0.0,
168.                                             trainable=True,
169.                                             dtype=tf.float32)
170.                 beta2 = tf.nn.softplus(soft_beta2)
171.                 log_beta2 = tf.log(tf.clip_by_value(beta2, 1e-8, 1e8))
172.         else:
173.             beta2 = tf.constant(fixed_cost_penalty)
174.             log_beta2 = tf.log(tf.clip_by_value(beta2, 1e-8, 1e8))
175.     else:

```

```

176.     beta2 = 0.0 # costs do not contribute to policy optimization
177.     print('Not using costs')
178.
179.     # Experience buffer
180.     replay_buffer = ReplayBuffer(obs_dim=obs_dim, act_dim=act_dim, size=replay_size, load_buffer=load_buffer)
181.
182.     # Count variables
183.     if proc_id() == 0:
184.         var_counts = tuple(count_vars(scope) for scope in
185.                             ['main/pi', 'main/qr1', 'main/qr2', 'main/qc1', 'main/qc2', 'main'])
186.         print((
187.             '\nNumber of parameters: \t pi: %d, \t qr1: %d, \t qr2: %d, \t qc1: %d, \t qc2: %d, \t total: %d\n') % var_co
188.             unts)
189.
190.     # Min Double-Q:
191.     min_q_pi = tf.minimum(qr1_pi, qr2_pi)
192.     min_q_pi_targ = tf.minimum(qr1_pi_targ, qr2_pi_targ)
193.
194.     # Targets for Q and V regression
195.     q_backup = tf.stop_gradient(r_ph + gamma * (1 - d_ph) * (min_q_pi_targ - alpha * logp_pi2))
196.     qc1_backup = tf.stop_gradient(c1_ph + gamma * (1 - d_ph) * qc1_pi_targ)
197.     qc2_backup = tf.stop_gradient(c2_ph + gamma * (1 - d_ph) * qc2_pi_targ)
198.
199.     cost_constraint = cost_lim * (1 - gamma ** max_ep_len) / (1 - gamma) / max_ep_len
200.     damp1 = damp_scale * tf.reduce_mean(cost_constraint - qc1)
201.     damp2 = damp_scale * tf.reduce_mean(cost_constraint - qc2)
202.
203.     # LBDSAC losses
204.     pi_loss = tf.reduce_mean(alpha * logp_pi - min_q_pi + (beta1 - damp1) * qc1_pi + (beta2 - damp2) * qc2_pi)
205.     qr1_loss = 0.5 * tf.reduce_mean((q_backup - qr1) ** 2)
206.     qr2_loss = 0.5 * tf.reduce_mean((q_backup - qr2) ** 2)
207.     qc1_loss = 0.5 * tf.reduce_mean((qc1_backup - qc1) ** 2)
208.     qc2_loss = 0.5 * tf.reduce_mean((qc2_backup - qc2) ** 2)
209.     q_loss = qr1_loss + qr2_loss + qc1_loss + qc2_loss
210.
211.     # Loss for alpha
212.     entropy_constraint *= act_dim
213.     pi_entropy = -tf.reduce_mean(logp_pi)
214.     alpha_loss = - alpha * (entropy_constraint - pi_entropy)
215.     print('using entropy constraint', entropy_constraint)
216.
217.     # Loss for beta
218.     if use_costs:
219.         if cost_constraint is None:
220.             cost_constraint = cost_lim * (1 - gamma ** max_ep_len) / (1 - gamma) / max_ep_len

```

```

219.     print('using cost constraint', cost_constraint)
220.     beta1_loss = beta1 * (cost_constraint - qc1)
221.     beta2_loss = beta2 * (cost_constraint - qc2)
222.
223.     # Policy train op
224.     train_pi_op = MpiAdamOptimizer(learning_rate=lr).minimize(pi_loss, var_list=get_vars('main/pi'), name='train_pi'
    )
225.
226.     # Value train op
227.     with tf.control_dependencies([train_pi_op]):
228.         train_q_op = MpiAdamOptimizer(learning_rate=lr).minimize(q_loss, var_list=get_vars('main/q'), name='train_q'
    )
229.
230.     if fixed_entropy_bonus is None:
231.         entreg_optimizer = MpiAdamOptimizer(learning_rate=lr)
232.         with tf.control_dependencies([train_q_op]):
233.             train_entreg_op = entreg_optimizer.minimize(alpha_loss, var_list=get_vars('entreg'))
234.
235.     if use_costs and fixed_cost_penalty is None:
236.         costpen_optimizer = MpiAdamOptimizer(learning_rate=lr * lr_scale)
237.
238.     if fixed_entropy_bonus is None:
239.         with tf.control_dependencies([train_entreg_op]):
240.             train_costpen1_op = costpen_optimizer.minimize(beta1_loss, var_list=get_vars('costpen1'))
241.     else:
242.         with tf.control_dependencies([train_q_op]):
243.             train_costpen1_op = costpen_optimizer.minimize(beta1_loss, var_list=get_vars('costpen1'))
244.
245.     if fixed_entropy_bonus is None:
246.         with tf.control_dependencies([train_entreg_op]):
247.             train_costpen2_op = costpen_optimizer.minimize(beta2_loss, var_list=get_vars('costpen2'))
248.     else:
249.         with tf.control_dependencies([train_q_op]):
250.             train_costpen2_op = costpen_optimizer.minimize(beta2_loss, var_list=get_vars('costpen2'))
251.
252.     target_update = get_target_update('main', 'target', polyak)
253.
254.     with tf.control_dependencies([train_pi_op]):
255.         with tf.control_dependencies([train_q_op]):
256.             grouped_update = tf.group([target_update])
257.
258.     if fixed_entropy_bonus is None:
259.         grouped_update = tf.group([grouped_update, train_entreg_op])
260.     if use_costs and fixed_cost_penalty is None:

```

```

261.     grouped_update = tf.group([grouped_update, train_costpen1_op])
262.     grouped_update = tf.group([grouped_update, train_costpen2_op])
263.
264.     def get_action(o, deterministic=False):
265.         act_op = mu if deterministic else pi
266.         return sess.run(act_op, feed_dict={x_ph: o.reshape(1, -1)})[0]
267.
268.     config = tf.ConfigProto()
269.     config.gpu_options.per_process_gpu_memory_fraction = 0.8
270.     sess = tf.Session(config=config)
271.
272.     writer = tf.summary.FileWriter(".logs/", sess.graph)
273.     # Initializing targets to match main variables
274.     target_init = get_target_update('main', 'target', 0.0)
275.     sess.run(tf.global_variables_initializer())
276.     sess.run(target_init)
277.     sess.run(sync_all_params())
278.     # Setup model saving
279.     logger.setup_tf_saver(sess, inputs={'x': x_ph, 'a': a_ph},
280.                           outputs={'mu': mu, 'pi': pi, 'qr1': qr1, 'qr2': qr2, 'qc1': qc1, 'qc2': qc2})
281.     start_time = time.time()
282.     o, r, d, ep_ret, ep_cost1, ep_cost2, ep_len= env.reset(), 0, False, 0, 0, 0
283.     total_steps = steps_per_epoch * epochs * max_ep_len
284.
285.     # variables to measure in an update
286.     vars_to_get = dict(LossPi=pi_loss, LossQR1=qr1_loss, LossQR2=qr2_loss, LossQC1=qc1_loss, LossQC2=qc2_lo
ss,
287.                       QR1Vals=qr1, QR2Vals=qr2, QC1Vals=qc1, QC2Vals=qc2, LogPi=logp_pi, PiEntropy=pi_entropy,
288.                       Alpha=alpha, LogAlpha=log_alpha, LossAlpha=alpha_loss)
289.     if use_costs:
290.         vars_to_get.update(dict(Beta1=beta1, LogBeta1=log_beta1, LossBeta1=beta1_loss,
291.                               Beta2=beta2, LogBeta2=log_beta2, LossBeta2=beta2_loss,))
292.
293.     print('starting training', proc_id())
294.     cum_cost1 = 0
295.     cum_cost2 = 0
296.     local_steps = 0
297.     local_steps_per_epoch = steps_per_epoch // num_procs()
298.     local_batch_size = batch_size // num_procs()
299.     epoch_start_time = time.time()
300.     for t in range(total_steps // num_procs()):
301.         if t > local_start_steps:
302.             a = get_action(o)
303.         else:

```

```

304.     a = env.action_sample()
305.     # Step the env
306.     o2, r, c1, c2, d = env.step(a,o,ep_len)
307.     r *= reward_scale # yee-haw
308.     ep_ret += r
309.     ep_cost1 += c1
310.     ep_cost2 += c2
311.     ep_len += 1
312.     local_steps += 1
313.     cum_cost1 += c1
314.     cum_cost2 += c2
315.     replay_buffer.store(o, a, r, o2, d, c1, c2)
316.     o = o2
317.     if d or (ep_len == max_ep_len):
318.         logger.store(EpRet=ep_ret, EpCost1=ep_cost1, EpCost2=ep_cost2)
319.         o, r, d, ep_ret, ep_cost1, ep_cost2, ep_len = env.reset(n = np.random.randint(n_wind_data)),\
320.             0, False, 0, 0, 0, 0
321.
322.     if t > 0 and t % update_freq == 0:
323.         for j in range(200):
324.             batch = replay_buffer.sample_batch(local_batch_size)
325.             feed_dict = {x_ph: batch['obs1'],
326.                 x2_ph: batch['obs2'],
327.                 a_ph: batch['acts'],
328.                 r_ph: batch['rews'],
329.                 c1_ph: batch['costs1'],
330.                 c2_ph: batch['costs2'],
331.                 d_ph: batch['done'],
332.                 }
333.             if t < local_update_after:
334.                 logger.store(**sess.run(vars_to_get, feed_dict))
335.             else:
336.                 values, _ = sess.run([vars_to_get, grouped_update], feed_dict)
337.                 logger.store(**values)
338.             ETA = (time.time()-epoch_start_time)*(total_steps//update_freq-t//update_freq)
339.             print(("Training: [epoch:%d|%d] [batch:%d|%d] ETA %d:%d:%d%\n"
340.                 % (t // (local_steps_per_epoch * max_ep_len),epochs-1,\
341.                     (t%(local_steps_per_epoch * max_ep_len)//max_ep_len),steps_per_epoch-1,\
342.                     ETA//3600, ETA%3600//60, ETA%60)))
343.             epoch_start_time = time.time()
344.
345.     # End of epoch wrap-up
346.     if t > 0 and t % (local_steps_per_epoch * max_ep_len) == 0:
347.         epoch = t // (local_steps_per_epoch * max_ep_len)

```



```
348.     cumulative_cost1 = mpi_sum(cum_cost1)
349.     cumulative_cost2 = mpi_sum(cum_cost2)
350.     cost_rate1 = cumulative_cost1 / ((epoch + 1) * steps_per_epoch)
351.     cost_rate2 = cumulative_cost2 / ((epoch + 1) * steps_per_epoch)
352.
353.     # Save model
354.     if (epoch % save_freq == 0) or (epoch == epochs - 1):
355.         saver = tf.train.Saver()
356.         saver.save(sess, './saved_models/Case118_%d/Model'%epoch)
357.         logger.log_tabular('Epoch', epoch)
358.         logger.log_tabular('EpRet', with_min_and_max=True)
359.         logger.log_tabular('EpCost1', with_min_and_max=True)
360.         logger.log_tabular('EpCost2', with_min_and_max=True)
361.         logger.log_tabular('CumulativeCost1', cumulative_cost1)
362.         logger.log_tabular('CumulativeCost2', cumulative_cost2)
363.         logger.log_tabular('CostRate1', cost_rate1)
364.         logger.log_tabular('CostRate2', cost_rate2)
365.         logger.log_tabular('LossPi', with_min_and_max=True)
366.         logger.log_tabular('LossQR1', with_min_and_max=True)
367.         logger.log_tabular('LossQC1', with_min_and_max=True)
368.         logger.log_tabular('LossQC2', with_min_and_max=True)
369.         logger.log_tabular('PiEntropy', with_min_and_max=True)
370.         logger.log_tabular('TotalTime', time.time() - start_time)
371.         logger.dump_tabular()
372.     writer.close()
```